

title: "Homework 1" author: "ArjunLaxman{style='background-color: yellow;}'" toc: true title-block-banner: true title-block-style: default format: pdf #format: pdf

[Link to the Github repository](#)

::: {.callout-important} ## Due: Fri, Jan 26, 2024 @ 11:59pm

```
install.packages("magrittr", repos = "https://cran.r-project.org")
```

The downloaded binary packages are in
/var/folders/kr/zyr76x5136x61sbwmb2_w_qr0000gn/T//RtmpnLHjLY/downloaded_packages

```
library(magrittr)
```

Please read the instructions carefully before submitting your assignment.

1. This assignment requires you to:

- Upload your Quarto markdown files to a [git](#) repository
- Upload a [PDF](#) file on Canvas

1. Don't collapse any code cells before submitting.

2. Remember to make sure all your code output is rendered properly before uploading your submission.

⚠ Please add your name to the the author information in the frontmatter before submitting your assignment. :::

Question 1

20 points

In this question, we will walk through the process of *forking* a [git](#) repository and submitting a *pull request*.

1. Navigate to the Github repository [here](#) and fork it by clicking on the icon in the top right



Provide a sensible name for your forked repository when prompted.

2. Clone your Github repository on your local machine

```
$ git clone <<https://github.com/arjunlaxman/hw1>>
$ cd hw-1
```

3. In order to activate the R environment for the homework, make sure you have `renv` installed beforehand. To activate the `renv` environment for this assignment, open an instance of the R console from within the directory and type

```
renv::activate()
```

Follow the instructions in order to make sure that `renv` is configured correctly.

4. Work on the *reminaing part* of this assignment as a `.qmd` file.

- Create a `PDF` and `HTML` file for your output by modifying the YAML frontmatter for the Quarto `.qmd` document

5. When you're done working on your assignment, push the changes to your github repository.

6. Navigate to the original Github repository [here](#) and submit a pull request linking to your repository.

Remember to **include your name** in the pull request information!

If you're stuck at any step along the way, you can refer to the [official Github docs here](#)

Question 2

30 points

Consider the following vector

```
my_vec <- c(
  "+0.07",
```

```
"-0.07",  
"+0.25",  
"-0.84",  
"+0.32",  
"-0.24",  
"-0.97",  
"-0.36",  
"+1.76",  
"-0.36"  
)
```

For the following questions, provide your answers in a code cell.

1. What data type does the vector contain?

```
data_type <- typeof(my_vec)
```

1. Create two new vectors called `my_vec_double` and `my_vec_int` which converts `my_vec` to Double & Integer types, respectively,

```
my_vec_double <- as.numeric(my_vec) # Converts character to double  
my_vec_int <- as.integer(my_vec_double) # Converts double to int
```

1. Create a new vector `my_vec_bool` which comprises of:

- TRUE if an element in `my_vec_double` is ≤ 0

```
my_vec_bool <- my_vec_double <= 0  
# TRUE if an element is <= 0  
# FALSE if an element is > 0
```

* ``FALSE`` if an element in `my_vec_double` is ≥ 0

How many elements of `my_vec_double` are greater than zero?

```
count_greater_than_zero <- sum(my_vec_double > 0)
```

1. Sort the values of `my_vec_double` in ascending order.

```
my_vec_sorted <- sort(my_vec_double)
```

Question 3

50 points

In this question we will get a better understanding of how R handles large data structures in memory.

1. Provide R code to construct the following matrices:

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} \quad \text{and} \quad \begin{bmatrix} 1 & 2 & 3 & 4 & 5 & \dots & 100 \\ 1 & 4 & 9 & 16 & 25 & \dots & 10000 \end{bmatrix}$$

Tip

```
# Creating the first 3x3 matrix
matrix_1 <- matrix(1:9, nrow=3, ncol=3, byrow=TRUE)

# Prints the matrix to check
print(matrix_1)
```

```
 [,1] [,2] [,3]
[1,]  1   2   3
[2,]  4   5   6
[3,]  7   8   9
```

```
# Creating the second matrix with integers and their squares
matrix_2 <- matrix(c(1:100, (1:100)^2), nrow=2, byrow=TRUE)

# Print the matrix to check
print(matrix_2)
```

```
 [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14]
[1,]  1   2   3   4   5   6   7   8   9   10   11   12   13   14
[2,]  1   4   9  16  25  36  49  64  81  100  121  144  169  196
 [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25] [,26]
[1,]  15  16  17  18  19  20  21  22  23  24  25  26
[2,] 225 256 289 324 361 400 441 484 529 576 625 676
 [,27] [,28] [,29] [,30] [,31] [,32] [,33] [,34] [,35] [,36] [,37] [,38]
[1,]  27  28  29  30  31  32  33  34  35  36  37  38
[2,] 729 784 841 900 961 1024 1089 1156 1225 1296 1369 1444
 [,39] [,40] [,41] [,42] [,43] [,44] [,45] [,46] [,47] [,48] [,49] [,50]
[1,]  39  40  41  42  43  44  45  46  47  48  49  50
[2,] 1521 1600 1681 1764 1849 1936 2025 2116 2209 2304 2401 2500
 [,51] [,52] [,53] [,54] [,55] [,56] [,57] [,58] [,59] [,60] [,61] [,62]
[1,]  51  52  53  54  55  56  57  58  59  60  61  62
[2,] 2601 2704 2809 2916 3025 3136 3249 3364 3481 3600 3721 3844
 [,63] [,64] [,65] [,66] [,67] [,68] [,69] [,70] [,71] [,72] [,73] [,74]
```

```

[1,] 63 64 65 66 67 68 69 70 71 72 73 74
[2,] 3969 4096 4225 4356 4489 4624 4761 4900 5041 5184 5329 5476
      [,75] [,76] [,77] [,78] [,79] [,80] [,81] [,82] [,83] [,84] [,85] [,86]
[1,] 75 76 77 78 79 80 81 82 83 84 85 86
[2,] 5625 5776 5929 6084 6241 6400 6561 6724 6889 7056 7225 7396
      [,87] [,88] [,89] [,90] [,91] [,92] [,93] [,94] [,95] [,96] [,97] [,98]
[1,] 87 88 89 90 91 92 93 94 95 96 97 98
[2,] 7569 7744 7921 8100 8281 8464 8649 8836 9025 9216 9409 9604
      [,99] [,100]
[1,] 99 100
[2,] 9801 10000

```

Recall the discussion in class on how R fills in matrices

In the next part, we will discover how knowledge of the way in which a matrix is stored in memory can inform better code choices. To this end, the following function takes an input n and creates an $n \times n$ matrix with random entries.

```

generate_matrix <- function(n){
  return(
    matrix(
      rnorm(n^2),
      nrow=n
    )
  )
}

```

For example:

```
generate_matrix(4)
```

```

      [,1]      [,2]      [,3]      [,4]
[1,] -0.8463432 -0.8555337 -0.2711772 -0.8549438
[2,] -0.1123393  2.1796750 -0.3168731  1.1889898
[3,] -0.1948057  0.8855451  0.1722450 -0.7926015
[4,]  0.1310771  0.2764002  0.3573066 -0.2179375

```

Let M be a fixed 50×50 matrix

```

M <- generate_matrix(50)
mean(M)

```

```
[1] -0.01957708
```

- Write a function `row_wise_scan` which scans the entries of M one row after another and outputs the number of elements whose value is ≥ 0 . You can use the following **starter code**

```

row_wise_scan <- function(x){
  n <- nrow(x) # Number of rows
  m <- ncol(x) # Number of columns

```

```

count <- 0
for(i in 1:n){
  for(j in 1:m){
    if(x[i, j] >= 0){
      count <- count + 1
    }
  }
}
return(count)
}

```

3. Similarly, write a function `col_wise_scan` which does exactly the same thing but scans the entries of `M` one column after another

```

col_wise_scan <- function(x){
  n <- nrow(x)
  m <- ncol(x)

  count <- 0

  for(j in 1:m){
    for(i in 1:n){
      if(x[i, j] >= 0){
        count <- count + 1
      }
    }
  }
  return(count)
}

```

You can check if your code is doing what it's supposed to using the function [here](#)¹

4. Between `col_wise_scan` and `row_wise_scan`, which function do you expect to take shorter to run? Why?

Since R stores matrices in column-major order, accessing data column-wise aligns with continuous memory locations, enhancing cache efficiency. Thus, the `col_wise_scan` function should perform faster than `row_wise_scan` due to more optimal memory access.

5. Write a function `time_scan` which takes in a method `f` and a matrix `M` and outputs the amount of time taken to run `f(M)`

```

time_scan <- function(f, M){
  initial_time <- Sys.time() # to Capture start time

  f(M) # Executes function f on matrix M

  final_time <- Sys.time() # Captures the end time

  total_time_taken <- final_time - initial_time # Calculates the duration
}

```

```
    return(total_time_taken)
}
```

Provide your output to

```
M <- generate_matrix(50) # Create a 50x50 matrix

# Measure the time taken by each scan function
row_wise_time <- time_scan(row_wise_scan, M)
col_wise_time <- time_scan(col_wise_scan, M)

# Output the times
list(
  row_wise_time = row_wise_time,
  col_wise_time = col_wise_time
)
```

```
$row_wise_time
Time difference of 0.002444983 secs
```

```
$col_wise_time
Time difference of 0.002273083 secs
```

Which took longer to run?

6. Repeat this experiment now when:

- M is a 100×100 matrix
- M is a 1000×1000 matrix
- M is a 5000×5000 matrix

```
# List of matrix sizes
sizes <- c(100, 1000, 5000)
results <- list()

# Loop through each size and measure performance
for (size in sizes) {
  M <- generate_matrix(size) # Generates matrix
  row_time <- time_scan(row_wise_scan, M) # Time taken row-wise scanning
  col_time <- time_scan(col_wise_scan, M) # Time taken column-wise scanning
  results[[paste(size, "x", size)]] <- list(row_time = row_time, col_time = col_time)
}

# Print results
results
```

```
$`100 x 100`
$`100 x 100`$row_time
Time difference of 0.0003459454 secs
```

```
$`100 x 100`$col_time  
Time difference of 0.0003318787 secs
```

```
$`1000 x 1000`  
$`1000 x 1000`$row_time  
Time difference of 0.0330019 secs
```

```
$`1000 x 1000`$col_time  
Time difference of 0.03213406 secs
```

```
$`5000 x 5000`  
$`5000 x 5000`$row_time  
Time difference of 1.069341 secs
```

```
$`5000 x 5000`$col_time  
Time difference of 0.8146741 secs
```

What can you conclude?

size of the matrix affects the performance of row-wise and column-wise scanning in R

Appendix

Print your R session information using the following command

```
sessionInfo()
```

```
R version 4.3.3 (2024-02-29)  
Platform: aarch64-apple-darwin20 (64-bit)  
Running under: macOS Sonoma 14.4.1
```

```
Matrix products: default  
BLAS: /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/lib/libRblas.0.dylib  
LAPACK: /Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/lib/libRlapack.dylib; LAPACK version 3.11.0
```

```
locale:  
[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
```

```
time zone: America/New_York  
tzcode source: internal
```

```
attached base packages:  
[1] stats      graphics  grDevices  utils      datasets  methods   base
```

```
other attached packages:  
[1] magrittr_2.0.3
```


loaded via a namespace (and not attached):

[1] compiler_4.3.3	fastmap_1.1.1	cli_3.6.2	tools_4.3.3
[5] htmltools_0.5.8.1	rmarkdown_2.26	knitr_1.45	jsonlite_1.8.8
[9] xfun_0.43	digest_0.6.35	rlang_1.1.3	evaluate_0.23

```
sapply(1:100, function(i) {  
  x <- generate_matrix(100)  
  row_wise_scan(x) == col_wise_scan(x)  
}) %>% sum == 100
```

[1] TRUE

Footnotes

1. If your code is right, the following code should evaluate to be TRUE 